

```
2020-02-15 02:01:33,094 INFO org.apache.hadoop.yarn.client.RMPProxy (main): Connecting to
ResourceManager at ip-172-31-65-239.ec2.internal/172.31.65.239:8032
2020-02-15 02:01:35,203 INFO org.apache.hadoop.mapreduce.lib.input.FileInputFormat (main):
Total input files to process : 1
2020-02-15 02:01:35,227 INFO com.hadoop.compression.lzo.GPLNativeCodeLoader (main): Loaded
native gpl library
2020-02-15 02:01:35,230 INFO com.hadoop.compression.lzo.LzoCodec (main): Successfully loaded
& initialized native-lzo library [hadoop-lzo rev 5f788d5e8f90539ee331702c753fa250727128f4]
2020-02-15 02:01:35,623 INFO org.apache.hadoop.mapreduce.JobSubmitter (main): number of
splits:20
2020-02-15 02:01:35,802 INFO org.apache.hadoop.mapreduce.JobSubmitter (main): Submitting
tokens for job: job_1581732032450_0001
2020-02-15 02:01:36,366 INFO org.apache.hadoop.yarn.client.api.impl.YarnClientImpl (main):
Submitted application application_1581732032450_0001
2020-02-15 02:01:36,472 INFO org.apache.hadoop.mapreduce.Job (main): The url to track the
job: http://ip-172-31-65-239.ec2.internal:20888/proxy/application_1581732032450_0001/
2020-02-15 02:01:36,473 INFO org.apache.hadoop.mapreduce.Job (main): Running job:
job_1581732032450_0001
2020-02-15 02:01:42,552 INFO org.apache.hadoop.mapreduce.Job (main): Job
job_1581732032450_0001 running in uber mode : false
2020-02-15 02:01:42,553 INFO org.apache.hadoop.mapreduce.Job (main): map 0% reduce 0%
2020-02-15 02:01:55,643 INFO org.apache.hadoop.mapreduce.Job (main): map 15% reduce 0%
2020-02-15 02:01:56,646 INFO org.apache.hadoop.mapreduce.Job (main): map 20% reduce 0%
2020-02-15 02:02:01,663 INFO org.apache.hadoop.mapreduce.Job (main): map 25% reduce 0%
2020-02-15 02:02:02,666 INFO org.apache.hadoop.mapreduce.Job (main): map 47% reduce 0%
2020-02-15 02:02:03,670 INFO org.apache.hadoop.mapreduce.Job (main): map 69% reduce 0%
2020-02-15 02:02:04,684 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 0%
2020-02-15 02:02:10,709 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 9%
2020-02-15 02:02:12,715 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 25%
2020-02-15 02:02:15,725 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 26%
2020-02-15 02:02:17,731 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 35%
2020-02-15 02:02:18,734 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 45%
2020-02-15 02:02:19,737 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 61%
2020-02-15 02:02:20,740 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 87%
2020-02-15 02:02:21,743 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 95%
2020-02-15 02:02:23,751 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 96%
2020-02-15 02:02:25,757 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 97%
2020-02-15 02:02:27,763 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 98%
2020-02-15 02:02:39,796 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 99%
2020-02-15 02:02:51,832 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 100%
2020-02-15 02:02:51,836 INFO org.apache.hadoop.mapreduce.Job (main): Job
job_1581732032450_0001 completed successfully
2020-02-15 02:02:51,937 INFO org.apache.hadoop.mapreduce.Job (main): Counters: 56
```

File System Counters

```
FILE: Number of bytes read=4787785
FILE: Number of bytes written=15209519
FILE: Number of read operations=0
FILE: Number of large read operations=0
FILE: Number of write operations=0
HDFS: Number of bytes read=2240
HDFS: Number of bytes written=4466136009
HDFS: Number of read operations=73
HDFS: Number of large read operations=0
HDFS: Number of write operations=22
S3: Number of bytes read=1319493757
S3: Number of bytes written=0
S3: Number of read operations=0
S3: Number of large read operations=0
S3: Number of write operations=0
```

Job Counters

```
Killed map tasks=1
Killed reduce tasks=1
Launched map tasks=20
Launched reduce tasks=12
```

Data-local map tasks=20
Total time spent by all maps in occupied slots (ms)=33220512
Total time spent by all reduces in occupied slots (ms)=43773504
Total time spent by all map tasks (ms)=346047
Total time spent by all reduce tasks (ms)=227987
Total vcore-milliseconds taken by all map tasks=346047
Total vcore-milliseconds taken by all reduce tasks=227987
Total megabyte-milliseconds taken by all map tasks=1063056384
Total megabyte-milliseconds taken by all reduce tasks=1400752128

Map-Reduce Framework

Map input records=85331845
Map output records=724126
Map output bytes=11700474
Map output materialized bytes=5156159
Input split bytes=2240
Combine input records=0
Combine output records=0
Reduce input groups=19999
Reduce shuffle bytes=5156159
Reduce input records=724126
Reduce output records=278819985
Spilled Records=1448252
Shuffled Maps =220
Failed Shuffles=0
Merged Map outputs=220
GC time elapsed (ms)=12394
CPU time spent (ms)=473880
Physical memory (bytes) snapshot=36593405952
Virtual memory (bytes) snapshot=173421760512
Total committed heap usage (bytes)=36207853568

Shuffle Errors

BAD_ID=0
CONNECTION=0
IO_ERROR=0
WRONG_LENGTH=0
WRONG_MAP=0
WRONG_REDUCE=0

File Input Format Counters

Bytes Read=1319493757

File Output Format Counters

Bytes Written=4466136009

2020-02-15 02:02:51,995 INFO org.apache.hadoop.yarn.client.RMPProxy (main): Connecting to ResourceManager at ip-172-31-65-239.ec2.internal/172.31.65.239:8032
2020-02-15 02:02:52,302 INFO org.apache.hadoop.mapreduce.lib.input.FileInputFormat (main): Total input files to process : 12
2020-02-15 02:02:52,349 INFO org.apache.hadoop.mapreduce.JobSubmitter (main): number of splits:57
2020-02-15 02:02:52,376 INFO org.apache.hadoop.mapreduce.JobSubmitter (main): Submitting tokens for job: job_1581732032450_0002
2020-02-15 02:02:52,405 INFO org.apache.hadoop.yarn.client.api.impl.YarnClientImpl (main): Submitted application application_1581732032450_0002
2020-02-15 02:02:52,418 INFO org.apache.hadoop.mapreduce.Job (main): The url to track the job: http://ip-172-31-65-239.ec2.internal:20888/proxy/application_1581732032450_0002/
2020-02-15 02:02:52,418 INFO org.apache.hadoop.mapreduce.Job (main): Running job: job_1581732032450_0002
2020-02-15 02:02:57,467 INFO org.apache.hadoop.mapreduce.Job (main): Job job_1581732032450_0002 running in uber mode : false
2020-02-15 02:02:57,467 INFO org.apache.hadoop.mapreduce.Job (main): map 0% reduce 0%
2020-02-15 02:03:09,533 INFO org.apache.hadoop.mapreduce.Job (main): map 2% reduce 0%
2020-02-15 02:03:10,536 INFO org.apache.hadoop.mapreduce.Job (main): map 4% reduce 0%
2020-02-15 02:03:14,549 INFO org.apache.hadoop.mapreduce.Job (main): map 5% reduce 0%
2020-02-15 02:03:16,555 INFO org.apache.hadoop.mapreduce.Job (main): map 9% reduce 0%
2020-02-15 02:03:17,560 INFO org.apache.hadoop.mapreduce.Job (main): map 30% reduce 0%
2020-02-15 02:03:18,563 INFO org.apache.hadoop.mapreduce.Job (main): map 36% reduce 0%
2020-02-15 02:03:25,587 INFO org.apache.hadoop.mapreduce.Job (main): map 37% reduce 0%

```
2020-02-15 02:03:26,590 INFO org.apache.hadoop.mapreduce.Job (main): map 38% reduce 0%
2020-02-15 02:03:32,611 INFO org.apache.hadoop.mapreduce.Job (main): map 39% reduce 0%
2020-02-15 02:03:33,614 INFO org.apache.hadoop.mapreduce.Job (main): map 40% reduce 0%
2020-02-15 02:03:34,617 INFO org.apache.hadoop.mapreduce.Job (main): map 41% reduce 0%
2020-02-15 02:03:35,620 INFO org.apache.hadoop.mapreduce.Job (main): map 44% reduce 0%
2020-02-15 02:03:36,623 INFO org.apache.hadoop.mapreduce.Job (main): map 46% reduce 1%
2020-02-15 02:03:39,632 INFO org.apache.hadoop.mapreduce.Job (main): map 47% reduce 1%
2020-02-15 02:03:40,638 INFO org.apache.hadoop.mapreduce.Job (main): map 48% reduce 1%
2020-02-15 02:03:41,641 INFO org.apache.hadoop.mapreduce.Job (main): map 49% reduce 1%
2020-02-15 02:03:42,644 INFO org.apache.hadoop.mapreduce.Job (main): map 51% reduce 1%
2020-02-15 02:03:45,652 INFO org.apache.hadoop.mapreduce.Job (main): map 52% reduce 1%
2020-02-15 02:03:47,658 INFO org.apache.hadoop.mapreduce.Job (main): map 53% reduce 1%
2020-02-15 02:03:50,666 INFO org.apache.hadoop.mapreduce.Job (main): map 54% reduce 1%
2020-02-15 02:03:52,671 INFO org.apache.hadoop.mapreduce.Job (main): map 55% reduce 1%
2020-02-15 02:03:53,673 INFO org.apache.hadoop.mapreduce.Job (main): map 56% reduce 1%
2020-02-15 02:03:54,676 INFO org.apache.hadoop.mapreduce.Job (main): map 59% reduce 1%
2020-02-15 02:03:59,688 INFO org.apache.hadoop.mapreduce.Job (main): map 60% reduce 1%
2020-02-15 02:04:01,693 INFO org.apache.hadoop.mapreduce.Job (main): map 60% reduce 3%
2020-02-15 02:04:02,696 INFO org.apache.hadoop.mapreduce.Job (main): map 61% reduce 3%
2020-02-15 02:04:05,702 INFO org.apache.hadoop.mapreduce.Job (main): map 62% reduce 3%
2020-02-15 02:04:06,704 INFO org.apache.hadoop.mapreduce.Job (main): map 63% reduce 3%
2020-02-15 02:04:08,709 INFO org.apache.hadoop.mapreduce.Job (main): map 64% reduce 3%
2020-02-15 02:04:10,718 INFO org.apache.hadoop.mapreduce.Job (main): map 64% reduce 4%
2020-02-15 02:04:11,720 INFO org.apache.hadoop.mapreduce.Job (main): map 65% reduce 4%
2020-02-15 02:04:12,722 INFO org.apache.hadoop.mapreduce.Job (main): map 67% reduce 4%
2020-02-15 02:04:14,727 INFO org.apache.hadoop.mapreduce.Job (main): map 68% reduce 4%
2020-02-15 02:04:16,731 INFO org.apache.hadoop.mapreduce.Job (main): map 68% reduce 5%
2020-02-15 02:04:17,734 INFO org.apache.hadoop.mapreduce.Job (main): map 69% reduce 5%
2020-02-15 02:04:18,737 INFO org.apache.hadoop.mapreduce.Job (main): map 70% reduce 7%
2020-02-15 02:04:19,739 INFO org.apache.hadoop.mapreduce.Job (main): map 70% reduce 8%
2020-02-15 02:04:24,751 INFO org.apache.hadoop.mapreduce.Job (main): map 70% reduce 11%
2020-02-15 02:04:28,761 INFO org.apache.hadoop.mapreduce.Job (main): map 71% reduce 11%
2020-02-15 02:04:31,768 INFO org.apache.hadoop.mapreduce.Job (main): map 72% reduce 11%
2020-02-15 02:04:32,770 INFO org.apache.hadoop.mapreduce.Job (main): map 73% reduce 11%
2020-02-15 02:04:36,779 INFO org.apache.hadoop.mapreduce.Job (main): map 75% reduce 11%
2020-02-15 02:04:38,783 INFO org.apache.hadoop.mapreduce.Job (main): map 77% reduce 11%
2020-02-15 02:04:42,791 INFO org.apache.hadoop.mapreduce.Job (main): map 78% reduce 11%
2020-02-15 02:04:44,797 INFO org.apache.hadoop.mapreduce.Job (main): map 79% reduce 11%
2020-02-15 02:04:46,801 INFO org.apache.hadoop.mapreduce.Job (main): map 80% reduce 11%
2020-02-15 02:04:49,807 INFO org.apache.hadoop.mapreduce.Job (main): map 81% reduce 11%
2020-02-15 02:04:52,814 INFO org.apache.hadoop.mapreduce.Job (main): map 82% reduce 11%
2020-02-15 02:04:54,818 INFO org.apache.hadoop.mapreduce.Job (main): map 83% reduce 11%
2020-02-15 02:04:56,822 INFO org.apache.hadoop.mapreduce.Job (main): map 84% reduce 11%
2020-02-15 02:05:02,834 INFO org.apache.hadoop.mapreduce.Job (main): map 85% reduce 11%
2020-02-15 02:05:04,839 INFO org.apache.hadoop.mapreduce.Job (main): map 86% reduce 11%
2020-02-15 02:05:06,843 INFO org.apache.hadoop.mapreduce.Job (main): map 87% reduce 11%
2020-02-15 02:05:07,846 INFO org.apache.hadoop.mapreduce.Job (main): map 88% reduce 11%
2020-02-15 02:05:08,848 INFO org.apache.hadoop.mapreduce.Job (main): map 89% reduce 11%
2020-02-15 02:05:10,852 INFO org.apache.hadoop.mapreduce.Job (main): map 90% reduce 11%
2020-02-15 02:05:12,856 INFO org.apache.hadoop.mapreduce.Job (main): map 91% reduce 12%
2020-02-15 02:05:13,858 INFO org.apache.hadoop.mapreduce.Job (main): map 91% reduce 13%
2020-02-15 02:05:18,869 INFO org.apache.hadoop.mapreduce.Job (main): map 91% reduce 14%
2020-02-15 02:05:20,874 INFO org.apache.hadoop.mapreduce.Job (main): map 92% reduce 14%
2020-02-15 02:05:22,879 INFO org.apache.hadoop.mapreduce.Job (main): map 94% reduce 14%
2020-02-15 02:05:23,882 INFO org.apache.hadoop.mapreduce.Job (main): map 96% reduce 14%
2020-02-15 02:05:25,886 INFO org.apache.hadoop.mapreduce.Job (main): map 97% reduce 14%
2020-02-15 02:05:26,887 INFO org.apache.hadoop.mapreduce.Job (main): map 98% reduce 14%
2020-02-15 02:05:28,891 INFO org.apache.hadoop.mapreduce.Job (main): map 99% reduce 20%
2020-02-15 02:05:29,893 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 20%
2020-02-15 02:05:31,897 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 21%
2020-02-15 02:05:34,903 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 22%
2020-02-15 02:05:36,907 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 26%
2020-02-15 02:05:37,909 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 31%
2020-02-15 02:05:40,914 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 35%
2020-02-15 02:05:41,916 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 40%
```

```
2020-02-15 02:05:42,919 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 44%
2020-02-15 02:05:43,921 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 51%
2020-02-15 02:05:46,934 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 60%
2020-02-15 02:05:47,936 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 61%
2020-02-15 02:05:48,938 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 63%
2020-02-15 02:05:49,940 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 66%
2020-02-15 02:05:52,950 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 68%
2020-02-15 02:05:53,952 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 69%
2020-02-15 02:05:54,953 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 70%
2020-02-15 02:05:55,955 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 72%
2020-02-15 02:05:58,961 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 74%
2020-02-15 02:05:59,963 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 75%
2020-02-15 02:06:00,965 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 76%
2020-02-15 02:06:01,967 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 78%
2020-02-15 02:06:04,973 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 80%
2020-02-15 02:06:05,977 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 81%
2020-02-15 02:06:06,978 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 82%
2020-02-15 02:06:07,980 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 84%
2020-02-15 02:06:10,986 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 86%
2020-02-15 02:06:11,988 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 87%
2020-02-15 02:06:12,990 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 88%
2020-02-15 02:06:13,992 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 90%
2020-02-15 02:06:16,997 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 92%
2020-02-15 02:06:17,999 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 93%
2020-02-15 02:06:19,001 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 94%
2020-02-15 02:06:20,008 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 96%
2020-02-15 02:06:23,014 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 98%
2020-02-15 02:06:25,018 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 99%
2020-02-15 02:06:26,020 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 100%
2020-02-15 02:06:29,028 INFO org.apache.hadoop.mapreduce.Job (main): Job
job_1581732032450_0002 completed successfully
2020-02-15 02:06:29,053 INFO org.apache.hadoop.mapreduce.Job (main): Counters: 56
```

File System Counters

```
FILE: Number of bytes read=7467363295
FILE: Number of bytes written=11158440979
FILE: Number of read operations=0
FILE: Number of large read operations=0
FILE: Number of write operations=0
HDFS: Number of bytes read=4467847402
HDFS: Number of bytes written=94052829
HDFS: Number of read operations=184
HDFS: Number of large read operations=0
HDFS: Number of write operations=22
S3: Number of bytes read=1319526325
S3: Number of bytes written=0
S3: Number of read operations=0
S3: Number of large read operations=0
S3: Number of write operations=0
```

Job Counters

```
Killed map tasks=1
Launched map tasks=57
Launched reduce tasks=11
Data-local map tasks=20
Rack-local map tasks=37
Total time spent by all maps in occupied slots (ms)=203527872
Total time spent by all reduces in occupied slots (ms)=213114240
Total time spent by all map tasks (ms)=2120082
Total time spent by all reduce tasks (ms)=1109970
Total vcore-milliseconds taken by all map tasks=2120082
Total vcore-milliseconds taken by all reduce tasks=1109970
Total megabyte-milliseconds taken by all map tasks=6512891904
Total megabyte-milliseconds taken by all reduce tasks=6819655680
```

Map-Reduce Framework

```
Map input records=364151830
Map output records=558364096
```

Map output bytes=11943550883
Map output materialized bytes=3724351672
Input split bytes=7457
Combine input records=0
Combine output records=0
Reduce input groups=19999
Reduce shuffle bytes=3724351672
Reduce input records=558364096
Reduce output records=7234833
Spilled Records=1670576088
Shuffled Maps =627
Failed Shuffles=0
Merged Map outputs=627
GC time elapsed (ms)=36866
CPU time spent (ms)=2515750
Physical memory (bytes) snapshot=96488894464
Virtual memory (bytes) snapshot=343878361088
Total committed heap usage (bytes)=99864281088

Shuffle Errors

BAD_ID=0
CONNECTION=0
IO_ERROR=0
WRONG_LENGTH=0
WRONG_MAP=0
WRONG_REDUCE=0

File Input Format Counters

Bytes Read=5787366270

File Output Format Counters

Bytes Written=94052829

2020-02-15 02:06:29,108 INFO org.apache.hadoop.yarn.client.RMProxy (main): Connecting to Resource Manager at ip-172-31-65-239.ec2.internal/172.31.65.239:8032
2020-02-15 02:06:29,228 INFO org.apache.hadoop.mapreduce.lib.input.FileInputFormat (main): Total input files to process : 11
2020-02-15 02:06:29,268 INFO org.apache.hadoop.mapreduce.JobSubmitter (main): number of splits:11
2020-02-15 02:06:29,288 INFO org.apache.hadoop.mapreduce.JobSubmitter (main): Submitting tokens for job: job_1581732032450_0003
2020-02-15 02:06:29,498 INFO org.apache.hadoop.yarn.client.api.impl.YarnClientImpl (main): Submitted application application_1581732032450_0003
2020-02-15 02:06:29,501 INFO org.apache.hadoop.mapreduce.Job (main): The url to track the job: http://ip-172-31-65-239.ec2.internal:20888/proxy/application_1581732032450_0003/
2020-02-15 02:06:29,501 INFO org.apache.hadoop.mapreduce.Job (main): Running job: job_1581732032450_0003
2020-02-15 02:06:35,546 INFO org.apache.hadoop.mapreduce.Job (main): Job job_1581732032450_0003 running in uber mode : false
2020-02-15 02:06:35,546 INFO org.apache.hadoop.mapreduce.Job (main): map 0% reduce 0%
2020-02-15 02:06:46,583 INFO org.apache.hadoop.mapreduce.Job (main): map 18% reduce 0%
2020-02-15 02:06:47,586 INFO org.apache.hadoop.mapreduce.Job (main): map 27% reduce 0%
2020-02-15 02:06:49,591 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 0%
2020-02-15 02:06:52,599 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 9%
2020-02-15 02:06:54,604 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 55%
2020-02-15 02:06:56,609 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 73%
2020-02-15 02:06:57,612 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 100%
2020-02-15 02:06:58,618 INFO org.apache.hadoop.mapreduce.Job (main): Job job_1581732032450_0003 completed successfully
2020-02-15 02:06:58,641 INFO org.apache.hadoop.mapreduce.Job (main): Counters: 56

File System Counters

FILE: Number of bytes read=5148714
FILE: Number of bytes written=14035019
FILE: Number of read operations=0
FILE: Number of large read operations=0
FILE: Number of write operations=0
HDFS: Number of bytes read=94054468
HDFS: Number of bytes written=0
HDFS: Number of read operations=22

```
HDFS: Number of large read operations=0
HDFS: Number of write operations=0
S3: Number of bytes read=0
S3: Number of bytes written=45
S3: Number of read operations=0
S3: Number of large read operations=0
S3: Number of write operations=0
Job Counters
  Killed map tasks=1
  Launched map tasks=11
  Launched reduce tasks=11
  Data-local map tasks=8
  Rack-local map tasks=3
  Total time spent by all maps in occupied slots (ms)=11524512
  Total time spent by all reduces in occupied slots (ms)=12925632
  Total time spent by all map tasks (ms)=120047
  Total time spent by all reduce tasks (ms)=67321
  Total vcore-milliseconds taken by all map tasks=120047
  Total vcore-milliseconds taken by all reduce tasks=67321
  Total megabyte-milliseconds taken by all map tasks=368784384
  Total megabyte-milliseconds taken by all reduce tasks=413620224
Map-Reduce Framework
  Map input records=7234833
  Map output records=7234833
  Map output bytes=94052829
  Map output materialized bytes=5150440
  Input split bytes=1639
  Combine input records=0
  Combine output records=0
  Reduce input groups=1
  Reduce shuffle bytes=5150440
  Reduce input records=7234833
  Reduce output records=1
  Spilled Records=14469666
  Shuffled Maps =121
  Failed Shuffles=0
  Merged Map outputs=121
  GC time elapsed (ms)=8095
  CPU time spent (ms)=56600
  Physical memory (bytes) snapshot=13933797376
  Virtual memory (bytes) snapshot=131637825536
  Total committed heap usage (bytes)=14030471168
Shuffle Errors
  BAD_ID=0
  CONNECTION=0
  IO_ERROR=0
  WRONG_LENGTH=0
  WRONG_MAP=0
  WRONG_REDUCE=0
File Input Format Counters
  Bytes Read=94052829
File Output Format Counters
  Bytes Written=45
```