

```
2020-02-15 06:21:12,538 INFO org.apache.hadoop.yarn.client.RMPProxy (main): Connecting to
ResourceManager at ip-172-31-70-67.ec2.internal/172.31.70.67:8032
2020-02-15 06:21:14,363 INFO org.apache.hadoop.mapreduce.lib.input.FileInputFormat (main):
Total input files to process : 1
2020-02-15 06:21:14,374 INFO com.hadoop.compression.lzo.GPLNativeCodeLoader (main): Loaded
native gpl library
2020-02-15 06:21:14,377 INFO com.hadoop.compression.lzo.LzoCodec (main): Successfully loaded
& initialized native-lzo library [hadoop-lzo rev 5f788d5e8f90539ee331702c753fa250727128f4]
2020-02-15 06:21:14,743 INFO org.apache.hadoop.mapreduce.JobSubmitter (main): number of
splits:20
2020-02-15 06:21:14,855 INFO org.apache.hadoop.mapreduce.JobSubmitter (main): Submitting
tokens for job: job_1581747592016_0001
2020-02-15 06:21:15,291 INFO org.apache.hadoop.yarn.client.api.impl.YarnClientImpl (main):
Submitted application application_1581747592016_0001
2020-02-15 06:21:15,342 INFO org.apache.hadoop.mapreduce.Job (main): The url to track the
job: http://ip-172-31-70-67.ec2.internal:20888/proxy/application_1581747592016_0001/
2020-02-15 06:21:15,342 INFO org.apache.hadoop.mapreduce.Job (main): Running job:
job_1581747592016_0001
2020-02-15 06:21:59,531 INFO org.apache.hadoop.mapreduce.Job (main): Job
job_1581747592016_0001 running in uber mode : false
2020-02-15 06:21:59,532 INFO org.apache.hadoop.mapreduce.Job (main): map 0% reduce 0%
2020-02-15 06:22:11,621 INFO org.apache.hadoop.mapreduce.Job (main): map 5% reduce 0%
2020-02-15 06:22:12,631 INFO org.apache.hadoop.mapreduce.Job (main): map 35% reduce 0%
2020-02-15 06:22:13,635 INFO org.apache.hadoop.mapreduce.Job (main): map 40% reduce 0%
2020-02-15 06:22:18,652 INFO org.apache.hadoop.mapreduce.Job (main): map 45% reduce 0%
2020-02-15 06:22:19,658 INFO org.apache.hadoop.mapreduce.Job (main): map 64% reduce 0%
2020-02-15 06:22:20,662 INFO org.apache.hadoop.mapreduce.Job (main): map 86% reduce 0%
2020-02-15 06:22:21,666 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 0%
2020-02-15 06:22:28,694 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 14%
2020-02-15 06:22:29,697 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 49%
2020-02-15 06:22:32,707 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 50%
2020-02-15 06:22:33,710 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 51%
2020-02-15 06:22:34,714 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 60%
2020-02-15 06:22:35,717 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 75%
2020-02-15 06:22:37,724 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 96%
2020-02-15 06:22:38,727 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 98%
2020-02-15 06:22:43,742 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 99%
2020-02-15 06:22:54,774 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 100%
2020-02-15 06:22:55,781 INFO org.apache.hadoop.mapreduce.Job (main): Job
job_1581747592016_0001 completed successfully
2020-02-15 06:22:55,869 INFO org.apache.hadoop.mapreduce.Job (main): Counters: 56
```

File System Counters

```
FILE: Number of bytes read=4784522
FILE: Number of bytes written=15540870
FILE: Number of read operations=0
FILE: Number of large read operations=0
FILE: Number of write operations=0
HDFS: Number of bytes read=2240
HDFS: Number of bytes written=4466136009
HDFS: Number of read operations=79
HDFS: Number of large read operations=0
HDFS: Number of write operations=26
S3: Number of bytes read=1319487717
S3: Number of bytes written=0
S3: Number of read operations=0
S3: Number of large read operations=0
S3: Number of write operations=0
```

Job Counters

```
Killed map tasks=1
Killed reduce tasks=1
Launched map tasks=20
Launched reduce tasks=14
Data-local map tasks=20
Total time spent by all maps in occupied slots (ms)=30064896
```

Total time spent by all reduces in occupied slots (ms)=49392384
Total time spent by all map tasks (ms)=313176
Total time spent by all reduce tasks (ms)=257252
Total vcore-milliseconds taken by all map tasks=313176
Total vcore-milliseconds taken by all reduce tasks=257252
Total megabyte-milliseconds taken by all map tasks=962076672
Total megabyte-milliseconds taken by all reduce tasks=1580556288

Map-Reduce Framework

Map input records=85331845
Map output records=724126
Map output bytes=11700474
Map output materialized bytes=5151070
Input split bytes=2240
Combine input records=0
Combine output records=0
Reduce input groups=19999
Reduce shuffle bytes=5151070
Reduce input records=724126
Reduce output records=278819985
Spilled Records=1448252
Shuffled Maps =260
Failed Shuffles=0
Merged Map outputs=260
GC time elapsed (ms)=11728
CPU time spent (ms)=473020
Physical memory (bytes) snapshot=37744017408
Virtual memory (bytes) snapshot=187907248128
Total committed heap usage (bytes)=38080086016

Shuffle Errors

BAD_ID=0
CONNECTION=0
IO_ERROR=0
WRONG_LENGTH=0
WRONG_MAP=0
WRONG_REDUCE=0

File Input Format Counters

Bytes Read=1319487717

File Output Format Counters

Bytes Written=4466136009

2020-02-15 06:22:55,915 INFO org.apache.hadoop.yarn.client.RMPProxy (main): Connecting to ResourceManager at ip-172-31-70-67.ec2.internal/172.31.70.67:8032
2020-02-15 06:22:56,170 INFO org.apache.hadoop.mapreduce.lib.input.FileInputFormat (main): Total input files to process : 14
2020-02-15 06:22:56,230 INFO org.apache.hadoop.mapreduce.JobSubmitter (main): number of splits:59
2020-02-15 06:22:56,267 INFO org.apache.hadoop.mapreduce.JobSubmitter (main): Submitting tokens for job: job_1581747592016_0002
2020-02-15 06:22:56,481 INFO org.apache.hadoop.yarn.client.api.impl.YarnClientImpl (main): Submitted application application_1581747592016_0002
2020-02-15 06:22:56,484 INFO org.apache.hadoop.mapreduce.Job (main): The url to track the job: http://ip-172-31-70-67.ec2.internal:20888/proxy/application_1581747592016_0002/
2020-02-15 06:22:56,484 INFO org.apache.hadoop.mapreduce.Job (main): Running job: job_1581747592016_0002
2020-02-15 06:23:01,533 INFO org.apache.hadoop.mapreduce.Job (main): Job job_1581747592016_0002 running in uber mode : false
2020-02-15 06:23:01,533 INFO org.apache.hadoop.mapreduce.Job (main): map 0% reduce 0%
2020-02-15 06:23:13,587 INFO org.apache.hadoop.mapreduce.Job (main): map 3% reduce 0%
2020-02-15 06:23:19,614 INFO org.apache.hadoop.mapreduce.Job (main): map 7% reduce 0%
2020-02-15 06:23:20,617 INFO org.apache.hadoop.mapreduce.Job (main): map 10% reduce 0%
2020-02-15 06:23:21,624 INFO org.apache.hadoop.mapreduce.Job (main): map 30% reduce 0%
2020-02-15 06:23:22,627 INFO org.apache.hadoop.mapreduce.Job (main): map 36% reduce 0%
2020-02-15 06:23:27,646 INFO org.apache.hadoop.mapreduce.Job (main): map 38% reduce 0%
2020-02-15 06:23:29,654 INFO org.apache.hadoop.mapreduce.Job (main): map 39% reduce 0%
2020-02-15 06:23:35,671 INFO org.apache.hadoop.mapreduce.Job (main): map 40% reduce 0%
2020-02-15 06:23:37,677 INFO org.apache.hadoop.mapreduce.Job (main): map 41% reduce 0%

```
2020-02-15 06:23:38,680 INFO org.apache.hadoop.mapreduce.Job (main): map 42% reduce 0%
2020-02-15 06:23:39,683 INFO org.apache.hadoop.mapreduce.Job (main): map 47% reduce 1%
2020-02-15 06:23:40,685 INFO org.apache.hadoop.mapreduce.Job (main): map 48% reduce 1%
2020-02-15 06:23:41,688 INFO org.apache.hadoop.mapreduce.Job (main): map 49% reduce 1%
2020-02-15 06:23:44,697 INFO org.apache.hadoop.mapreduce.Job (main): map 50% reduce 1%
2020-02-15 06:23:45,700 INFO org.apache.hadoop.mapreduce.Job (main): map 52% reduce 1%
2020-02-15 06:23:46,702 INFO org.apache.hadoop.mapreduce.Job (main): map 53% reduce 1%
2020-02-15 06:23:47,705 INFO org.apache.hadoop.mapreduce.Job (main): map 54% reduce 1%
2020-02-15 06:23:49,710 INFO org.apache.hadoop.mapreduce.Job (main): map 55% reduce 1%
2020-02-15 06:23:51,715 INFO org.apache.hadoop.mapreduce.Job (main): map 56% reduce 1%
2020-02-15 06:23:55,724 INFO org.apache.hadoop.mapreduce.Job (main): map 57% reduce 1%
2020-02-15 06:23:57,728 INFO org.apache.hadoop.mapreduce.Job (main): map 61% reduce 1%
2020-02-15 06:23:58,730 INFO org.apache.hadoop.mapreduce.Job (main): map 62% reduce 1%
2020-02-15 06:24:00,735 INFO org.apache.hadoop.mapreduce.Job (main): map 63% reduce 1%
2020-02-15 06:24:02,741 INFO org.apache.hadoop.mapreduce.Job (main): map 64% reduce 1%
2020-02-15 06:24:03,743 INFO org.apache.hadoop.mapreduce.Job (main): map 65% reduce 1%
2020-02-15 06:24:04,746 INFO org.apache.hadoop.mapreduce.Job (main): map 65% reduce 2%
2020-02-15 06:24:05,749 INFO org.apache.hadoop.mapreduce.Job (main): map 66% reduce 2%
2020-02-15 06:24:09,758 INFO org.apache.hadoop.mapreduce.Job (main): map 67% reduce 3%
2020-02-15 06:24:11,763 INFO org.apache.hadoop.mapreduce.Job (main): map 68% reduce 3%
2020-02-15 06:24:14,770 INFO org.apache.hadoop.mapreduce.Job (main): map 69% reduce 3%
2020-02-15 06:24:15,772 INFO org.apache.hadoop.mapreduce.Job (main): map 71% reduce 4%
2020-02-15 06:24:17,776 INFO org.apache.hadoop.mapreduce.Job (main): map 72% reduce 6%
2020-02-15 06:24:18,779 INFO org.apache.hadoop.mapreduce.Job (main): map 73% reduce 6%
2020-02-15 06:24:19,781 INFO org.apache.hadoop.mapreduce.Job (main): map 73% reduce 7%
2020-02-15 06:24:20,784 INFO org.apache.hadoop.mapreduce.Job (main): map 74% reduce 7%
2020-02-15 06:24:21,786 INFO org.apache.hadoop.mapreduce.Job (main): map 75% reduce 8%
2020-02-15 06:24:23,791 INFO org.apache.hadoop.mapreduce.Job (main): map 75% reduce 11%
2020-02-15 06:24:29,805 INFO org.apache.hadoop.mapreduce.Job (main): map 76% reduce 11%
2020-02-15 06:24:35,820 INFO org.apache.hadoop.mapreduce.Job (main): map 77% reduce 11%
2020-02-15 06:24:36,822 INFO org.apache.hadoop.mapreduce.Job (main): map 78% reduce 11%
2020-02-15 06:24:37,824 INFO org.apache.hadoop.mapreduce.Job (main): map 79% reduce 11%
2020-02-15 06:24:38,826 INFO org.apache.hadoop.mapreduce.Job (main): map 86% reduce 11%
2020-02-15 06:24:39,828 INFO org.apache.hadoop.mapreduce.Job (main): map 86% reduce 12%
2020-02-15 06:24:41,833 INFO org.apache.hadoop.mapreduce.Job (main): map 87% reduce 12%
2020-02-15 06:24:42,835 INFO org.apache.hadoop.mapreduce.Job (main): map 88% reduce 12%
2020-02-15 06:24:43,838 INFO org.apache.hadoop.mapreduce.Job (main): map 89% reduce 12%
2020-02-15 06:24:44,840 INFO org.apache.hadoop.mapreduce.Job (main): map 90% reduce 12%
2020-02-15 06:24:45,842 INFO org.apache.hadoop.mapreduce.Job (main): map 91% reduce 12%
2020-02-15 06:24:48,849 INFO org.apache.hadoop.mapreduce.Job (main): map 92% reduce 12%
2020-02-15 06:24:50,853 INFO org.apache.hadoop.mapreduce.Job (main): map 93% reduce 12%
2020-02-15 06:24:51,855 INFO org.apache.hadoop.mapreduce.Job (main): map 93% reduce 13%
2020-02-15 06:24:52,858 INFO org.apache.hadoop.mapreduce.Job (main): map 95% reduce 13%
2020-02-15 06:24:54,862 INFO org.apache.hadoop.mapreduce.Job (main): map 96% reduce 13%
2020-02-15 06:24:56,866 INFO org.apache.hadoop.mapreduce.Job (main): map 97% reduce 13%
2020-02-15 06:24:58,870 INFO org.apache.hadoop.mapreduce.Job (main): map 97% reduce 14%
2020-02-15 06:24:59,872 INFO org.apache.hadoop.mapreduce.Job (main): map 98% reduce 14%
2020-02-15 06:25:00,874 INFO org.apache.hadoop.mapreduce.Job (main): map 98% reduce 16%
2020-02-15 06:25:02,878 INFO org.apache.hadoop.mapreduce.Job (main): map 98% reduce 19%
2020-02-15 06:25:04,882 INFO org.apache.hadoop.mapreduce.Job (main): map 99% reduce 19%
2020-02-15 06:25:06,886 INFO org.apache.hadoop.mapreduce.Job (main): map 99% reduce 20%
2020-02-15 06:25:07,888 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 20%
2020-02-15 06:25:08,893 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 22%
2020-02-15 06:25:11,899 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 26%
2020-02-15 06:25:12,901 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 27%
2020-02-15 06:25:13,903 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 28%
2020-02-15 06:25:14,912 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 30%
2020-02-15 06:25:15,914 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 33%
2020-02-15 06:25:16,916 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 38%
2020-02-15 06:25:17,918 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 43%
2020-02-15 06:25:18,920 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 44%
2020-02-15 06:25:19,922 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 50%
2020-02-15 06:25:20,924 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 52%
2020-02-15 06:25:21,926 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 55%
2020-02-15 06:25:22,928 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 57%
```

```
2020-02-15 06:25:23,930 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 60%
2020-02-15 06:25:24,932 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 66%
2020-02-15 06:25:26,936 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 67%
2020-02-15 06:25:27,938 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 68%
2020-02-15 06:25:28,940 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 69%
2020-02-15 06:25:29,942 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 71%
2020-02-15 06:25:30,944 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 73%
2020-02-15 06:25:32,947 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 74%
2020-02-15 06:25:33,949 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 76%
2020-02-15 06:25:34,951 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 77%
2020-02-15 06:25:35,953 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 78%
2020-02-15 06:25:36,955 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 80%
2020-02-15 06:25:38,959 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 82%
2020-02-15 06:25:39,961 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 83%
2020-02-15 06:25:40,963 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 84%
2020-02-15 06:25:41,965 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 85%
2020-02-15 06:25:42,966 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 87%
2020-02-15 06:25:44,971 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 88%
2020-02-15 06:25:45,973 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 89%
2020-02-15 06:25:46,974 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 91%
2020-02-15 06:25:47,976 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 92%
2020-02-15 06:25:48,978 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 94%
2020-02-15 06:25:49,980 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 95%
2020-02-15 06:25:50,982 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 96%
2020-02-15 06:25:51,985 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 97%
2020-02-15 06:25:52,987 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 98%
2020-02-15 06:25:53,988 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 99%
2020-02-15 06:25:55,992 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 100%
2020-02-15 06:25:59,001 INFO org.apache.hadoop.mapreduce.Job (main): Job
job_1581747592016_0002 completed successfully
2020-02-15 06:25:59,023 INFO org.apache.hadoop.mapreduce.Job (main): Counters: 56
```

File System Counters

```
FILE: Number of bytes read=7439048606
FILE: Number of bytes written=11120550462
FILE: Number of read operations=0
FILE: Number of large read operations=0
FILE: Number of write operations=0
HDFS: Number of bytes read=4467847645
HDFS: Number of bytes written=94052829
HDFS: Number of read operations=196
HDFS: Number of large read operations=0
HDFS: Number of write operations=26
S3: Number of bytes read=1319505594
S3: Number of bytes written=0
S3: Number of read operations=0
S3: Number of large read operations=0
S3: Number of write operations=0
```

Job Counters

```
Killed map tasks=1
Launched map tasks=59
Launched reduce tasks=13
Data-local map tasks=20
Rack-local map tasks=39
Total time spent by all maps in occupied slots (ms)=203494560
Total time spent by all reduces in occupied slots (ms)=208605888
Total time spent by all map tasks (ms)=2119735
Total time spent by all reduce tasks (ms)=1086489
Total vcore-milliseconds taken by all map tasks=2119735
Total vcore-milliseconds taken by all reduce tasks=1086489
Total megabyte-milliseconds taken by all map tasks=6511825920
Total megabyte-milliseconds taken by all reduce tasks=6675388416
```

Map-Reduce Framework

```
Map input records=364151830
Map output records=558364096
Map output bytes=11943550883
```

Map output materialized bytes=3723045221
Input split bytes=7700
Combine input records=0
Combine output records=0
Reduce input groups=19999
Reduce shuffle bytes=3723045221
Reduce input records=558364096
Reduce output records=7234833
Spilled Records=1665316920
Shuffled Maps =767
Failed Shuffles=0
Merged Map outputs=767
GC time elapsed (ms)=39156
CPU time spent (ms)=2499260
Physical memory (bytes) snapshot=97176858624
Virtual memory (bytes) snapshot=367831834624
Total committed heap usage (bytes)=96383533056

Shuffle Errors

BAD_ID=0
CONNECTION=0
IO_ERROR=0
WRONG_LENGTH=0
WRONG_MAP=0
WRONG_REDUCE=0

File Input Format Counters

Bytes Read=5787345539

File Output Format Counters

Bytes Written=94052829

2020-02-15 06:25:59,066 INFO org.apache.hadoop.yarn.client.RMProxy (main): Connecting to ResourceManager at ip-172-31-70-67.ec2.internal/172.31.70.67:8032
2020-02-15 06:25:59,186 INFO org.apache.hadoop.mapreduce.lib.input.FileInputFormat (main): Total input files to process : 13
2020-02-15 06:25:59,234 INFO org.apache.hadoop.mapreduce.JobSubmitter (main): number of splits:13
2020-02-15 06:25:59,253 INFO org.apache.hadoop.mapreduce.JobSubmitter (main): Submitting tokens for job: job_1581747592016_0003
2020-02-15 06:25:59,468 INFO org.apache.hadoop.yarn.client.api.impl.YarnClientImpl (main): Submitted application application_1581747592016_0003
2020-02-15 06:25:59,471 INFO org.apache.hadoop.mapreduce.Job (main): The url to track the job: http://ip-172-31-70-67.ec2.internal:20888/proxy/application_1581747592016_0003/
2020-02-15 06:25:59,471 INFO org.apache.hadoop.mapreduce.Job (main): Running job: job_1581747592016_0003
2020-02-15 06:26:05,517 INFO org.apache.hadoop.mapreduce.Job (main): Job job_1581747592016_0003 running in uber mode : false
2020-02-15 06:26:05,517 INFO org.apache.hadoop.mapreduce.Job (main): map 0% reduce 0%
2020-02-15 06:26:13,562 INFO org.apache.hadoop.mapreduce.Job (main): map 8% reduce 0%
2020-02-15 06:26:14,565 INFO org.apache.hadoop.mapreduce.Job (main): map 15% reduce 0%
2020-02-15 06:26:16,571 INFO org.apache.hadoop.mapreduce.Job (main): map 38% reduce 0%
2020-02-15 06:26:18,577 INFO org.apache.hadoop.mapreduce.Job (main): map 62% reduce 0%
2020-02-15 06:26:19,582 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 0%
2020-02-15 06:26:21,588 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 8%
2020-02-15 06:26:22,591 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 54%
2020-02-15 06:26:23,594 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 62%
2020-02-15 06:26:24,597 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 77%
2020-02-15 06:26:26,601 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 92%
2020-02-15 06:26:29,608 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 100%
2020-02-15 06:26:30,614 INFO org.apache.hadoop.mapreduce.Job (main): Job job_1581747592016_0003 completed successfully
2020-02-15 06:26:30,637 INFO org.apache.hadoop.mapreduce.Job (main): Counters: 56

File System Counters

FILE: Number of bytes read=5148754
FILE: Number of bytes written=14715242
FILE: Number of read operations=0
FILE: Number of large read operations=0
FILE: Number of write operations=0

HDFS: Number of bytes read=94054753
HDFS: Number of bytes written=0
HDFS: Number of read operations=26
HDFS: Number of large read operations=0
HDFS: Number of write operations=0
S3: Number of bytes read=0
S3: Number of bytes written=45
S3: Number of read operations=0
S3: Number of large read operations=0
S3: Number of write operations=0

Job Counters

Killed map tasks=1
Launched map tasks=13
Launched reduce tasks=13
Data-local map tasks=11
Rack-local map tasks=2
Total time spent by all maps in occupied slots (ms)=12790080
Total time spent by all reduces in occupied slots (ms)=14845632
Total time spent by all map tasks (ms)=133230
Total time spent by all reduce tasks (ms)=77321
Total vcore-milliseconds taken by all map tasks=133230
Total vcore-milliseconds taken by all reduce tasks=77321
Total megabyte-milliseconds taken by all map tasks=409282560
Total megabyte-milliseconds taken by all reduce tasks=475060224

Map-Reduce Framework

Map input records=7234833
Map output records=7234833
Map output bytes=94052829
Map output materialized bytes=5151212
Input split bytes=1924
Combine input records=0
Combine output records=0
Reduce input groups=1
Reduce shuffle bytes=5151212
Reduce input records=7234833
Reduce output records=1
Spilled Records=14469666
Shuffled Maps =169
Failed Shuffles=0
Merged Map outputs=169
GC time elapsed (ms)=10511
CPU time spent (ms)=61050
Physical memory (bytes) snapshot=16662654976
Virtual memory (bytes) snapshot=155570987008
Total committed heap usage (bytes)=16934502400

Shuffle Errors

BAD_ID=0
CONNECTION=0
IO_ERROR=0
WRONG_LENGTH=0
WRONG_MAP=0
WRONG_REDUCE=0

File Input Format Counters

Bytes Read=94052829

File Output Format Counters

Bytes Written=45