

```
2020-02-15 04:20:23,809 INFO org.apache.hadoop.yarn.client.RMPProxy (main): Connecting to
ResourceManager at ip-172-31-64-193.ec2.internal/172.31.64.193:8032
2020-02-15 04:20:26,927 INFO org.apache.hadoop.mapreduce.lib.input.FileInputFormat (main):
Total input files to process : 1
2020-02-15 04:20:26,938 INFO com.hadoop.compression.lzo.GPLNativeCodeLoader (main): Loaded
native gpl library
2020-02-15 04:20:26,941 INFO com.hadoop.compression.lzo.LzoCodec (main): Successfully loaded
& initialized native-lzo library [hadoop-lzo rev 5f788d5e8f90539ee331702c753fa250727128f4]
2020-02-15 04:20:27,318 INFO org.apache.hadoop.mapreduce.JobSubmitter (main): number of
splits:20
2020-02-15 04:20:27,419 INFO org.apache.hadoop.mapreduce.JobSubmitter (main): Submitting
tokens for job: job_1581740342678_0001
2020-02-15 04:20:27,890 INFO org.apache.hadoop.yarn.client.api.impl.YarnClientImpl (main):
Submitted application application_1581740342678_0001
2020-02-15 04:20:27,956 INFO org.apache.hadoop.mapreduce.Job (main): The url to track the
job: http://ip-172-31-64-193.ec2.internal:20888/proxy/application_1581740342678_0001/
2020-02-15 04:20:27,957 INFO org.apache.hadoop.mapreduce.Job (main): Running job:
job_1581740342678_0001
2020-02-15 04:20:35,070 INFO org.apache.hadoop.mapreduce.Job (main): Job
job_1581740342678_0001 running in uber mode : false
2020-02-15 04:20:35,071 INFO org.apache.hadoop.mapreduce.Job (main): map 0% reduce 0%
2020-02-15 04:22:01,432 INFO org.apache.hadoop.mapreduce.Job (main): map 10% reduce 0%
2020-02-15 04:22:02,435 INFO org.apache.hadoop.mapreduce.Job (main): map 15% reduce 0%
2020-02-15 04:22:07,449 INFO org.apache.hadoop.mapreduce.Job (main): map 20% reduce 0%
2020-02-15 04:22:10,458 INFO org.apache.hadoop.mapreduce.Job (main): map 25% reduce 0%
2020-02-15 04:22:11,461 INFO org.apache.hadoop.mapreduce.Job (main): map 30% reduce 0%
2020-02-15 04:22:12,463 INFO org.apache.hadoop.mapreduce.Job (main): map 35% reduce 0%
2020-02-15 04:22:15,471 INFO org.apache.hadoop.mapreduce.Job (main): map 40% reduce 0%
2020-02-15 04:22:18,479 INFO org.apache.hadoop.mapreduce.Job (main): map 40% reduce 4%
2020-02-15 04:22:22,488 INFO org.apache.hadoop.mapreduce.Job (main): map 40% reduce 5%
2020-02-15 04:22:25,495 INFO org.apache.hadoop.mapreduce.Job (main): map 40% reduce 6%
2020-02-15 04:22:30,507 INFO org.apache.hadoop.mapreduce.Job (main): map 40% reduce 7%
2020-02-15 04:23:10,599 INFO org.apache.hadoop.mapreduce.Job (main): map 45% reduce 7%
2020-02-15 04:23:12,603 INFO org.apache.hadoop.mapreduce.Job (main): map 50% reduce 8%
2020-02-15 04:23:16,613 INFO org.apache.hadoop.mapreduce.Job (main): map 55% reduce 8%
2020-02-15 04:23:17,616 INFO org.apache.hadoop.mapreduce.Job (main): map 60% reduce 8%
2020-02-15 04:23:18,619 INFO org.apache.hadoop.mapreduce.Job (main): map 70% reduce 10%
2020-02-15 04:23:19,622 INFO org.apache.hadoop.mapreduce.Job (main): map 75% reduce 11%
2020-02-15 04:23:20,624 INFO org.apache.hadoop.mapreduce.Job (main): map 80% reduce 11%
2020-02-15 04:23:22,628 INFO org.apache.hadoop.mapreduce.Job (main): map 85% reduce 11%
2020-02-15 04:23:23,630 INFO org.apache.hadoop.mapreduce.Job (main): map 90% reduce 11%
2020-02-15 04:23:24,633 INFO org.apache.hadoop.mapreduce.Job (main): map 90% reduce 15%
2020-02-15 04:23:25,636 INFO org.apache.hadoop.mapreduce.Job (main): map 95% reduce 16%
2020-02-15 04:23:28,642 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 16%
2020-02-15 04:23:29,645 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 61%
2020-02-15 04:23:30,648 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 92%
2020-02-15 04:23:34,657 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 100%
2020-02-15 04:23:35,662 INFO org.apache.hadoop.mapreduce.Job (main): Job
job_1581740342678_0001 completed successfully
2020-02-15 04:23:35,770 INFO org.apache.hadoop.mapreduce.Job (main): Counters: 55
File System Counters
FILE: Number of bytes read=2077281
FILE: Number of bytes written=9767721
FILE: Number of read operations=0
FILE: Number of large read operations=0
FILE: Number of write operations=0
HDFS: Number of bytes read=2240
HDFS: Number of bytes written=0
HDFS: Number of read operations=20
HDFS: Number of large read operations=0
HDFS: Number of write operations=0
S3: Number of bytes read=27705142809
S3: Number of bytes written=8828733
S3: Number of read operations=0
```

S3: Number of large read operations=0

S3: Number of write operations=0

Job Counters

Killed map tasks=2

Launched map tasks=20

Launched reduce tasks=13

Data-local map tasks=20

Total time spent by all maps in occupied slots (ms)=255477408

Total time spent by all reduces in occupied slots (ms)=119723520

Total time spent by all map tasks (ms)=2661223

Total time spent by all reduce tasks (ms)=623560

Total vcore-milliseconds taken by all map tasks=2661223

Total vcore-milliseconds taken by all reduce tasks=623560

Total megabyte-milliseconds taken by all map tasks=8175277056

Total megabyte-milliseconds taken by all reduce tasks=3831152640

Map-Reduce Framework

Map input records=85331845

Map output records=400016

Map output bytes=8828733

Map output materialized bytes=2079027

Input split bytes=2240

Combine input records=0

Combine output records=0

Reduce input groups=1

Reduce shuffle bytes=2079027

Reduce input records=400016

Reduce output records=400016

Spilled Records=800032

Shuffled Maps =260

Failed Shuffles=0

Merged Map outputs=260

GC time elapsed (ms)=50373

CPU time spent (ms)=2532630

Physical memory (bytes) snapshot=33548099584

Virtual memory (bytes) snapshot=188015235072

Total committed heap usage (bytes)=32017219584

Shuffle Errors

BAD_ID=0

CONNECTION=0

IO_ERROR=0

WRONG_LENGTH=0

WRONG_MAP=0

WRONG_REDUCE=0

File Input Format Counters

Bytes Read=1319506309

File Output Format Counters

Bytes Written=8828733