

Assignment -2Report

PROBLEM ANALYSIS:

Cardinality:

Explanation:

Find the number of incoming and outgoing edges for every node. This way we can determine if Y has m incoming edges—from followers x_1, \dots, x_m —and n outgoing edges—to followed users z_1, \dots, z_n —then there exactly $m \cdot n$ length-2 paths through Y —one for each combination (x_i, Y, z_j) . This implies that if we can efficiently count the incoming and outgoing edges for each user. Now we can take the sum of this count to determine the cardinality.

PseudoCode:

```
class CardinalityMapper
    method Map(doc edge)
        for all users  $\in$  edge  $e$  do
            user = users.split(",")
            from = user[0]
            to = user[1]

            If from < MAX_FILTER && to < MAX_FILTER
                Emit(from, "O")
                Emit(to, "I")
```

```
class CardinalityReducer
    method Reduce(user, counts [O, O, . . . I, I, O])
        For list of O and I for a user:
            If O:
                countO++
            Else:
                countI++
        emit(user, countO*countI)
```

Class ComMapper

```

method Map(doc intermediate)
  for all count ∈ intermediate e do
    c = count.split(",")
    cnt = c[1]
    Emit("Sum", cnt)

```

Class ComReducer

```

method Reduce(Sum , counts [1, 3, . 320. .])
  Sum <- 0
  for all cnt ∈ intermediate e do
    Sum = Sum+cnt
  Emit("Cardinality ", Sum)

```

MAX_VALUE	CARDINALITY
1000	400016
10000	73597234
1000000	104405319023
11316811	953138453592

	RS join input	RS join shuffled	RS join output	Rep join input	Rep join file cache	Rep join output
Step 1 (join of Edges with itself)	950 B~	950 B~	950 B~	950 B~	950 B~	950 B~
Step 2 (join of Path2 with Edges)	950 B~	950 B~	950 B~	N/A	N/A	N/A

JOIN IMPLEMENTATION:

MAX FILTER

Explanation:

Read the file and in the map function emit only those edges where from and to values are less than the MAX_FILTER.

Pseudo-Code:

```
class MaxFilterMapper
    method Map(doc edge)
        for all users  $\in$  edge e do
            user = users.split(",")
            from = user[0]
            to = user[1]

            If from < MAX_FILTER && to < MAX_FILTER
                Emit(from, to)

class MaxFilterReducer
    method Reduce(from , counts [u1, u2, . . . , un])
        Emit(from, to)
```

REDUCE SIDE JOIN

Explanation:

Input:

```
1,2
2,3
3,1
```

First Phase:

In the first Phase of the mapper we will split the user graph and emit the from user as key and (from user, to user) as value and to user as key and (from user and to user) as value.

Mapper output:

```
1: (1,2)
1: (3,1)
2: (2,3)
```

2: (1,2)
3: (2,3)
3: (3,1)

In the first Reducer, for every from user create a hash table or lookup with from user as key and a list of to users as values.

So the hash map would look like for 1's iteration:

1 [2], 3[1]

Now we go through the list of users for various keys in the hashmap and see if we can reach the neighbouring node from that node.

Here, we will get no path for 2.

But for 1 we get:

3 -> 1 -> 2.

Now, lets see what happens when the key is 2.

Hashmap: 2[3] 1[2]

The result: 1 -> 2 -> 3

Similarly, when the key is 3.

Hashmap: 3[1] 2[3]

The result: 2 -> 3 -> 1

Thus, we get 3 paths from the reducer.

3 -> 1 -> 2
1 -> 2 -> 3
2 -> 3 -> 1

Once we got all the Two- paths. We can get the original edges and the new Two path and combine them in the mapper.

So the output of the second mapper would look something like:

1 : (1 -> 2 -> 3)
1: (3 -> 1)
1: (2 -> 3 -> 1)

1: (1 -> 2)

2: (2 -> 3 -> 1)

2: (2 -> 3)

2: (3 -> 1 -> 2)

2: (1 -> 2)

3: (3 -> 1 -> 2)

3: (3 -> 1)

3: (2 -> 3)

3: (1 -> 2 -> 3)

In reducer we make two hashmap

Threepath = 1[3]

Twopath = 3[1]

Now we iterate through the keys and try to see if the from key is present for the user in Twopath hashmap.

As we can see in the above example we have.

So we increase the triangle count by 1.

Similarly for 2's iteration

We will get

Threepath = 2[1]

Twopath = 1[2]

We increase the count by 1

For Similarly for 3's iteration

We will get

Threepath = 3[2]

Twopath = 2[3]

We increase the count by 1

So we have a total count of triangle as 3.

Since we have counted the same triangle thrice we will divide the result by 3 to obtain the number of triangle.

Pseudo-Code:

```
class PathTwoMapper
  method Map(doc edge)
    for all users  $\in$  edge e do
      user = users.split(",")
      from = user[0]
      to = user[1]
      If from < MAX_FILTER && to < MAX_FILTER
        Emit(from, to)

class PathTwoReducer
  method Reduce(from , counts [u1, u2, . . . , un])
    Hashmap hm <- (from, List(to))
    For all key from  $\in$  hashmap hm :
      For all neighbour  $\in$  hm(key):
        For all neighbors_neighbour  $\in$  hm(neighbour):
          Emit(key, neighbour, neighbors_neighbour)

//Read the original File as well as the output of the PathTwoReducer.
class TriangleMapper
  method Map(doc originalPath_TwoPath_combined)
    for all users  $\in$  originalPath_TwoPath_combined e do
      user = users.split(",")
      from = user[0]
      to = user[user.length -1]
      If from < MAX_FILTER && to < MAX_FILTER
        Emit(from, value)
        Emit(to, value)

class TriangleReducer
  method Reduce(from , counts [u1, u2, . . . , un])
    Hashmap threePath<- (from, List(to)) if user is from
    Hashmap twoPath<- (from, List(to)) if user is to
    For all key from  $\in$  hashmap threePath:
      For all neighbors_neighbour  $\in$  hm(key):
        If hashMap twoPath(neighbors_neighbour) contains key:
          Emit("Triangle", "1")
```

Class CounterMapper

Method map(doc triangles)

word1 = triangles.split(",")[0]

word2 = triangles.split(",")[1]

emit(word1, word2)

Class CounterReducer

Method reduce(Triangle , counts [1, 1, . . . , 1])

sum \leftarrow 0

for all count $c \in$ counts [1, 1, . . . , 1] do

sum \leftarrow sum + c

Emit("the number of triangles are", (count sum)/3)

REPLICATED JOIN:

Explanation:

First we will go through the edge file and create a adjacency list of every user.

Then in the mapper we will start reading the edge and try to find the "to users" neighbour's neighbour and see if it is equal to the from user. If it is equal to the from user, then we will increase the count of the triangle by 1. Since we count the same triangle thrice , we will divide the count by 3 to obtain the result.

Pseudo-Code:

HashMap adjList //Distributed cache

method ReplicatedJoinMap(doc edge)

for all users \in edge e do

user = users.split(",")

from = user[0]

to = user[1]

For neighbors_neighbor \in adjList(to):

If adj(neighbors_neighbor) contains from:

Emit("Triangle", "1")

Class ReplicatedJoinReducer

Method reduce(Triangle , counts [1, 1, . . . , 1])

sum \leftarrow 0

for all count c \in counts [1, 1, . . . , 1] do

sum \leftarrow sum + c

Emit("the number of triangles are", (count sum)/3)

Configuration	Small Cluster Result	Large Cluster Result
RS-join, MAX = 1000	Running time: 98 sec Triangle count: 1099	Running time: 100 sec Triangle count: 1099
Rep-join, MAX = 1000	Running time: 190 sec Triangle count: 1099	Running time: 196 sec Triangle count: 1099
RS-join, MAX = 10000	Running time: 170 sec Triangle count: 520296	Running time: 164 sec Triangle count: 520296
Rep-join, MAX = 10000	Running time: 196 sec Triangle count: 520296	Running time: 198 sec Triangle count: 520296
RS-join, MAX = 20000	Running time: 330 sec Triangle count: 2411611	Running time: 324 sec Triangle count: 2411611
Rep-join, MAX = 20000	Running time: 430 sec Triangle count: 2411611	Running time: 320 sec Triangle count: 2411611

